


SCORE: 80%

DATE TAKEN: Feb 14, 2020

**CONTEXT:** I purchased a set of 800+ exam questions from Udemi. It is divided into 6 exams. This is the first of the 6. The exam consisted of 160 questions and it was timed to a duration of 3.5 hours. I was able to complete this exam in just under 2.5 hours.

## Practice Test 1 : 160 Questions - Results

160 questions | 3 hours 40 minutes | 65% correct required to pass



## Practice Test 1 : 160 Questions - Results

[Return to review](#)

### Attempt 1

Incorrect

Question 15: **Incorrect**

**Which of the following is an example of synchronous replication which occurs in the AWS service?**



A. AWS RDS Read Replica's for MySQL, MariaDB and PostgreSQL

**(Incorrect)**



B. AWS Multi-AZ RDS

**(Correct)**



C. Redis engine for Amazon ElastiCache replication



D. AWS RDS Read Replica's for Oracle

### Explanation

Amazon RDS Multi-AZ deployments provide enhanced availability and durability for Database (DB) Instances, making them a natural fit for production database workloads. When you provision a Multi-AZ DB Instance, Amazon RDS automatically creates a primary DB Instance and synchronously replicates the data to a standby instance in a different Availability Zone (AZ). For more information on Multi-AZ, please visit the below URL:

<https://aws.amazon.com/rds/details/multi-az/> Option A is invalid because Amazon RDS takes a snapshot of the source instance and creates a read-only instance from the snapshot. For MySQL, MariaDB and PostgreSQL, Amazon RDS uses those engines' native asynchronous replication to update the read replica whenever there is a change to the source DB instance. Option C is invalid, because the Redis engine for Amazon ElastiCache supports replication with automatic failover, but the Redis engine's replication is asynchronous Option D is invalid because this is not supported by AWS.

Question 16: **Incorrect**

You want to get the reason for your EC2 Instance termination from the CLI. Which of the below commands is ideal in getting the reason.

☐

A. `aws ec2 describe-instances`

(Correct)

☐

B. `aws ec2 describe-images`

☐

C. `aws ec2 get-console-screenshot`

☒

D. `aws ec2 describe-volume-status`

(Incorrect)

### Explanation

When you execute the AWS `ec2 describe-instances` CLI command with the `instance_id` as shown below `aws ec2 describe-instances --instance-id instance_id` In the JSON response that's displayed, locate the `StateReason` element. An example is shown below. This will help in understanding why the instance was shutdown. `"StateReason": { "Message": "Client.UserInitiatedShutdown: User initiated shutdown", "Code": "Client.UserInitiatedShutdown" }`, Option B is invalid because this command describes one or more of the images (AMIs, AKIs, and ARIs) available to you Option C is invalid because retrieve a JPG-format screenshot of a running instance. This might not help to the complete extent of understanding why the instance was terminated. Option D is invalid because this command describes the status of the specified volumes. For more information on the command, please visit the below URL: <http://docs.aws.amazon.com/cli/latest/reference/ec2/describe-instances.html>

Question 19: **Incorrect**

**You have created your own VPC and subnet in AWS. You have launched an instance in that subnet. You have noticed that the instance is not receiving a DNS name. Which of the below options could be a valid reason for this issue.**

• ☐

A. The CIDR block for the VPC is invalid

• ☐

B. The CIDR block for the subnet is invalid

• ☒

C. The VPC configuration needs to be changed.

**(Correct)**

• ☐

D. The subnet configuration needs to be changed.

**(Incorrect)**

### **Explanation**

If the DNS hostnames option of the VPC is not set to 'Yes' then the instances launched in the subnet will not get DNS Names. You can change the option by choosing your VPC and clicking on 'Edit DNS Hostnames' Option A and B are invalid because if the CIDR blocks were invalid then the VPC or subnet would not be created. Option D is invalid because the subnet configuration does not have the effect on the DNS hostnames. For more information on VPC's, please visit the below URL: <https://aws.amazon.com/vpc/>

Question 24: **Incorrect**

**You are creating a number of EBS Volumes for your EC2 instances. You are concerned on the backups of the EBS Volumes. Which of the below is a way to backup the EBS Volumes**

• ☐

A. Configure Amazon Storage Gateway with EBS volumes as the data source and store the backups on premise through the storage gateway

• ☐

B. Write a cronjob that uses the AWS CLI to take a snapshot of production EBS volumes.

**(Correct)**

• ☒

C. Use a lifecycle policy to back up EBS volumes stored on Amazon S3 for durability

**(Incorrect)**

• ☐

D. Write a cronjob on the server that compresses the data and then copy it to Glacier

### **Explanation**

A point-in-time snapshot of an EBS volume, can be used as a baseline for new volumes or for data backup. If you make periodic snapshots of a volume, the snapshots are incremental—only the blocks on the device that have changed after your last snapshot are saved in the new snapshot. Even though snapshots are saved incrementally, the snapshot deletion process is designed so that you need to retain only the most recent snapshot in order to restore the entire volume. You can create a snapshot via the CLI command – create-snapshot Option A is incorrect because you normally use the Storage gateway to backup your on-premise data. Option C is incorrect because this is used for S3 storage Option D is incorrect because compression is another maintenance task and storing it in Glacier is not an ideal option For more information on snapshots, please visit the below URL:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-creating-snapshot.html>

Question 26: **Incorrect**

**You are trying to configure Cross Region Replication for your S3 bucket. But you are not able to select the option of Cross Region Replication and is disabled. Which of the below could be the possible reasons for this ?**

☐ A. The feature is not available in that region  
**(Incorrect)**

☒ B. You need to enable versioning on the bucket  
**(Correct)**

☐ C. The source region is currently down

☐ D. The destination region is currently down

### **Explanation**

Requirements for cross-region replication: The source and destination buckets must be versioning-enabled. The source and destination buckets must be in different AWS regions. You can replicate objects from a source bucket to only one destination bucket. Amazon S3 must have permission to replicate objects from that source bucket to the destination bucket on your behalf. If the source bucket owner also owns the object, the bucket owner has full permissions to replicate the object. If not, the source bucket owner must have permission for the Amazon S3 actions `s3:GetObjectVersion` and `s3:GetObjectVersionACL` to read the object and object ACL. If you are setting up cross-region replication in a cross-account scenario (where the source and destination buckets are owned by different AWS accounts), the source bucket owner must have permission to replicate objects in the destination bucket. The destination bucket owner needs to grant these permissions via a bucket policy. Option A is invalid , because it is available in all regions Option C is invalid because if so, then you would not be able to access S3 in that region Option D is invalid because you have not reached the configuration stage to select the destination bucket For more information on S3 Cross Region Replication, please visit the below URL:  
<https://docs.aws.amazon.com/AmazonS3/latest/dev/crr.html>

Question 31: **Incorrect**

**You want to ensure that you keep a check on the Active Volumes , Active snapshots and Elastic IP addresses you use so that you don't go beyond the service limit. Which of the below services can help in this regard?**



A. AWS Cloudwatch

**(Incorrect)**



B. AWS EC2



C. AWS Trusted Advisor

**(Correct)**



D. AWS SNS

### **Explanation**

An online resource to help you reduce cost, increase performance, and improve security by optimizing your AWS environment, Trusted Advisor provides real time guidance to help you provision your resources following AWS best practices. Below is a snapshot of the service limits it can monitor Option A is invalid because even though you can monitor resources , it cannot be checked against the service limit. Option B is invalid because this is the Elastic Compute cloud service Option D is invalid because it can be send notification but not check on service limits For more information on the Trusted Advisor monitoring, please visit the below URL:  
<https://aws.amazon.com/premiumsupport/ta-faqs/>

Question 35: **Incorrect**

**You have several AWS reserved instances in your account. They have been running for some time, but now need to be shutdown since they are no longer required. The data is still required for future purposes. Which of the below possible 2 steps can be taken.**

• ☐

A. Convert the instance to on-demand instances

• ☐

B. Sell the instances on the AWS Reserved Instance Marketplace

(Correct)

• ☒

C. Take snapshots of the EBS volumes and terminate the instances

(Correct)

• ☐

D. Convert the instance to spot instances

### Explanation

The Reserved Instance Marketplace is a platform that supports the sale of third-party and AWS customers' unused Standard Reserved Instances, which vary in term lengths and pricing options. For example, you may want to sell Reserved Instances after moving instances to a new AWS region, changing to a new instance type, ending projects before the term expiration, when your business needs change, or if you have unneeded capacity. For more information on selling instances, please visit the below URL:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ri-market-general.html> Since the data is still required, it's better to take snapshots of the existing volumes and then terminate the instances. For more information on EBS Snapshots, please visit the below URL:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html> Option A and D are invalid, because you cannot convert Reserved instances to either on-demand instances or Spot Instances.



Question 42: **Incorrect**

**You have some EC2 instances hosted in your AWS environment. You have a concern that not all of the EC2 instances are being utilized. Which of the below mentioned services can help you find underutilized resources in AWS ? Select 2 options.**

• ☒

A. AWS Cloudwatch

(Correct)

• ☐

B. SNS

• ☐

C. AWS Trusted Advisor

(Correct)

• ☒

D. Cloudtrail

(Incorrect)

### Explanation

The AWS Trusted Advisor can help you identify underutilized resources in AWS. For more information on AWS trusted advisor please visit the below URL:

<https://aws.amazon.com/premiumsupport/trustedadvisor/> If You look at the Cloudwatch graphs, the CPU utilization of your resources and you can see the trend over time in the graphs. For more information on AWS Cloudwatch please visit the below URL:

<https://aws.amazon.com/cloudwatch/>

Question 49: **Incorrect**

**In the shared responsibility model, what is the customer not responsible for?**

• ☐

A. Edge locations

(Correct)

• ☐

B. Installation of custom firewall software

• ☐

C. Security Groups

• ☒

D. Applying an SSL Certificate to an ELB

(Incorrect)

### Explanation

AWS has published the Shared Responsibility Model. And the Physical networking comes as part of the responsibility of AWS. For more information on the Shared Responsibility Model, please refer to the below URL: <https://aws.amazon.com/compliance/shared-responsibility-model/>

Question 65: **Incorrect**

**You are creating a Provisioned IOPS volume in AWS. The size of the volume is 8 GiB. Which of the following are the possible values that can put for the IOPS of the volume**

☐

A. 400

**(Correct)**

☐

B. 500

☐

C. 600

☒

D. 1000

**(Incorrect)**

### **Explanation**

The Maximum ratio of IOPS to volume size is 50:1 , so if the volume size is 8 GiB , the maximum IOPS of the volume can be 400. If you go beyond this value , you will get an error as shown in the screenshot below. For more information on Provisioned IOPS, please visit the below URL: <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSVolumeTypes.html>

Question 70: **Incorrect**

**Which of the below resources cannot be tagged in AWS**

• ☐

A. Images

• ☒

B. EBS Volumes

**(Incorrect)**

• ☐

C. VPC endpoint

**(Correct)**

• ☐

D. VPC

### **Explanation**

Tags enable you to categorize your AWS resources in different ways, for example, by purpose, owner, or environment. This is useful when you have many resources of the same type — you can quickly identify a specific resource based on the tags you've assigned to it. Each tag consists of a key and an optional value, both of which you define. But you cannot tag a VPC endpoint. For more information on AWS Resourcing Tagging, please visit [http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Using\\_Tags.html](http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Using_Tags.html)

Question 71: **Incorrect**

**What type of monitoring for EBS volumes is available automatically in 5 minute periods at no charge?**

☐

A. Basic

**(Correct)**

☐

B. Primary

☐

C. Detailed

☒

D. Local

**(Incorrect)**

**Explanation**

Visit the AWS documentation for the types of monitoring data. For more information on Volume monitoring, please visit <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/monitoring-volume-status.html>

Question 78: **Incorrect**

**After migrating an application architecture from on-premise to AWS you will not be responsible for the ongoing maintenance of packages for which of the following AWS services that your application uses. Choose the 2 correct answers from the options below.**

- ☒

A. Elastic Beanstalk  
(Incorrect)
- ☒

B. RDS  
(Correct)
- ☐

C. DynamoDB  
(Correct)
- ☐

D. EC2

### **Explanation**

Both RDS and DynamoDB are managed solutions provided by AWS. Amazon Relational Database Service (Amazon RDS) makes it easy to set up, operate, and scale a relational database in the cloud. It provides cost-efficient and resizable capacity while managing time-consuming database administration tasks, freeing you up to focus on your applications and business. For more information on RDS, please refer to the below link <https://aws.amazon.com/rds/> Amazon DynamoDB is a fast and flexible NoSQL database service for all applications that need consistent, single-digit millisecond latency at any scale. It is a fully managed cloud database and supports both document and key-value store models. For more information on DynamoDB, please refer to the below link <https://aws.amazon.com/dynamodb/>

Question 79: **Incorrect**

**What is the difference between an availability zone and an edge location? Choose the correct answer from the options below**

☐

A. Edge locations are used as control stations for AWS resources

☐

B. An edge location is used as a link when building load balancing between regions

☒

C. An Availability Zone is an isolated location inside a region; an edge location will deliver cached content to the closest location to reduce latency

**(Correct)**

☐

D. An availability zone is a grouping of AWS resources in a specific region; an edge location is a specific resource within the AWS region

**(Incorrect)**

### **Explanation**

Edge locations Using a network of edge locations around the world, Amazon CloudFront caches copies of your static content close to viewers, lowering latency when they download your objects and giving you the high, sustained data transfer rates needed to deliver large popular objects to end users at scale. For more information on Cloudfront and edge locations, please refer to the below link [https://aws.amazon.com/cloudfront/ Availability Zones](https://aws.amazon.com/cloudfront/AvailabilityZones) Each region is completely independent. Each Availability Zone is isolated, but the Availability Zones in a region are connected through low-latency links. The following diagram illustrates the relationship between regions and Availability Zones. For more information on AZ, please refer to the below link <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-regions-availability-zones.html>

Question 81: **Incorrect**

**Your supervisor asks you to create a decoupled application whose process includes dependencies on EC2 instances and servers located in your company's on-premises data center. Which of these are you least likely to recommend as part of that process? Choose the correct answer from the options below:**

• ☐

A. SQS polling from an EC2 instance deployed with an IAM role

• ☒

B. An SWF workflow

**(Incorrect)**

• ☐

C. SQS polling from an EC2 instance using IAM user credentials

**(Correct)**

• ☐

D. SQS polling from an on-premises server using IAM user credentials

### **Explanation**

Note that the question asks you for the least likely recommended option. The correct answer is C, SQS polling from an EC2 instance using IAM user credentials. An EC2 role should be used when deploying EC2 instances to grant permissions rather than storing IAM user credentials in EC2 instances. You should use IAM roles for secure communication between EC2 instances and resources on AWS. Your most likely scenario will actually be SQS polling from an EC2 instance deployed with an IAM role because when your polling SQS from EC2 you should use IAM roles. What you should never do is use IAM user api keys for authentication to poll sqs messages. An IAM role is similar to a user, in that it is an AWS identity with permission policies that determine what the identity can and cannot do in AWS. However, instead of being uniquely associated with one person, a role is intended to be assumable by anyone who needs it. Also, a role does not have any credentials (password or access keys) associated with it. Instead, if a user is assigned to a role, access keys are created dynamically and provided to the user. The most likely scenario will actually be SQS polling from an EC2 instance deployed with an IAM role because when your polling SQS from EC2 you should use IAM roles. We should never use IAM user api keys for authentication to poll SQS messages. Option C is correct which is least likely scenario. For more information on IAM Roles, please refer to the below link:  
[http://docs.aws.amazon.com/IAM/latest/UserGuide/id\\_roles.html](http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles.html)



Question 90: **Incorrect**

**Which of the following is incorrect with regards to Private IP addresses?**



A. In Amazon EC2 classic, the private IP addresses are only returned to Amazon EC2 when the instance is stopped or terminated

**(Incorrect)**



B. In Amazon VPC, an instance retains its private IP addresses when the instance is stopped.



C. In Amazon VPC, an instance does NOT retain its private IP addresses when the instance is stopped.

**(Correct)**



D. In Amazon EC2 classic, the private IP address is associated exclusively with the instance for its lifetime

### **Explanation**

The following is true with regards to Private IP addressing For instances launched in a VPC, a private IPv4 address remains associated with the network interface when the instance is stopped and restarted, and is released when the instance is terminated. For instances launched in EC2-Classic, we release the private IPv4 address when the instance is stopped or terminated. If you restart your stopped instance, it receives a new private IPv4 address For more information on IP addressing , please refer to the below link:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-instance-addressing.html>

Question 91: **Incorrect**

**Which of the following are best practices for monitoring your EC2 Instances**

- ☐ A. Create and implement a monitoring plan that collects monitoring data from all of the parts in your AWS solution

- ☒ B. Automate monitoring tasks as much as possible  
**(Incorrect)**

- ☐ C. Check the log files on your EC2 instances

- ☐ D. All of the above  
**(Correct)**

### **Explanation**

Use the following best practices for monitoring to help you with your Amazon EC2 monitoring tasks. Make monitoring a priority to head off small problems before they become big ones. Create and implement a monitoring plan that collects monitoring data from all of the parts in your AWS solution so that you can more easily debug a multi-point failure if one occurs. Your monitoring plan should address, at a minimum, the following questions: What are your goals for monitoring? What resources you will monitor? How often you will monitor these resources? What monitoring tools will you use? Who will perform the monitoring tasks? Who should be notified when something goes wrong? Automate monitoring tasks as much as possible. Check the log files on your EC2 instances. For more information on monitoring EC2 , please refer to the below link: [http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/monitoring\\_ec2.html](http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/monitoring_ec2.html)

Question 92: **Incorrect**

**You work for a major news network in Europe. They have just released a new app which allows users to report on events as and when they happen using their mobile phone. Users are able to upload pictures from the app and then other users will be able to view these pics. Your organization expects this app to grow very quickly, essentially doubling it's user base every month. The app uses S3 to store the media and you are expecting sudden and large increases in traffic to S3 when a major news event takes place as people will be uploading content in huge numbers). You need to keep your storage costs to a minimum however and it does not matter if some objects are lost. Which storage media should you use to keep costs as low as possible?**

☒ A. S3 – Infrequently Accessed Storage.

**(Incorrect)**

☐ B. S3 – Reduced Redundancy Storage (RRS).

**(Correct)**

☐ C. Glacier.

☐ D. S3 – Provisioned IOPS.

### **Explanation**

Since the requirement mentions that it does not matter if objects are lost and you need a low cost storage option then Reduced Redundancy Storage is the best option. The AWS Documentation mentions the below on Reduced Redundancy Storage Reduced Redundancy Storage (RRS) is an Amazon S3 storage option that enables customers to store noncritical, reproducible data at lower levels of redundancy than Amazon S3's standard storage. It provides a highly available solution for distributing or sharing content that is durably stored elsewhere, or for storing thumbnails, transcoded media, or other processed data that can be easily reproduced For more information on RRS, please refer to the below link: <https://aws.amazon.com/s3/reduced-redundancy/>

Question 96: **Incorrect**

**Is it true that EBS can always tolerate an Availability Zone failure?**

☐

A. No, all EBS volume is stored in a single Availability Zone

(Correct)

☒

B. Yes, EBS volume has multiple copies so it should be fine

(Incorrect)

☐

C. Depends on how it is setup

☐

D. Depends on the Region where EBS volume is initiated

### Explanation

EBS Volume replicated to physical hardware with in the same available zone, So if AZ fails then EBS volume will fail. That's why AWS recommend to always keep EBS volume snapshot in S3 bucket for high durability. "When you create an EBS volume in an Availability Zone, it is automatically replicated within that zone to prevent data loss due to the failure of any single hardware component." Option B is wrong as EBS volume has multiple copies but with in same AZ , so volume will not persist in case of AZ failure. Option C is wrong because there is no special setup available to persist EBS volume across region or AZ. Answer D is wrong as EBS volume has same behavior regardless of region. As per AWS user guide: <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSVolumes.html>

Question 101: **Incorrect**

A company wants to host a selection of MongoDB instances. They are expecting a high load and want to have as low latency as possible. Which class of instances from the below list should they choose from.

☐

A. T2

☐

B. I2

(Correct)

☐

C. T1

☒

D. G2

(Incorrect)

### Explanation

I2 instances are optimized to deliver tens of thousands of low-latency, random I/O operations per second (IOPS) to applications. They are well suited for the following scenarios: NoSQL databases (for example, Cassandra and MongoDB) Clustered databases Online transaction processing (OLTP) systems For more information on I2 instances, please visit the link <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/i2-instances.html>

Question 105: **Incorrect**

**For DynamoDB, what are the scenario's in which you would want to enable Cross-region replication?**

- ☐ A. Live data migration
- ☐ B. Easier Traffic management

• ☒ C. Disaster Recovery

**(Incorrect)**

• ☐ D. All of the above

**(Correct)**

### **Explanation**

From the AWS Documentation, it clearly states the reason for why you would want to enable Cross-Region Replication For more information on DynamoDB, please visit the url <https://aws.amazon.com/dynamodb/faqs/>

Question 106: **Incorrect**

**You have launched two web servers in private subnet and one ELB (internet facing) in public subnet in your VPC. Yet, you are still unable to access your web application through the internet, which of the following would likely the cause of this? Choose two correct options**

• ☒

A. Web server must be launched inside public subnet and not private subnet.

**(Incorrect)**

• ☐

B. Route table for public subnet is not configured to route to VPC internet gateway.

**(Correct)**

• ☐

C. No elastic IP is assigned to web servers.

• ☐

D. No internet gateway is attached to the VPC.

**(Correct)**

### **Explanation**

In order for the EC2 or ELBs to be accessible from internet, we would need to configure the route table for public subnet to route traffic to VPC internet gateway. For example: ? For information on VPC Route Tables and VPC Internet Gateway, please visit the link:  
[https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Route\\_Tables.html](https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Route_Tables.html)  
[https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Internet\\_Gateway.html](https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Internet_Gateway.html)

Question 109: **Incorrect**

**Which of the below instances is used normally for massive parallel computations?**

• ☐

A. Spot Instances

(Correct)

• ☐

B. On-Demand Instances

• ☒

C. Dedicated Instances

(Incorrect)

• ☐

D. This is not possible in AWS

### Explanation

This is clearly given in the aws documentation For more information on Spot Instances, please visit the link – <https://aws.amazon.com/ec2/spot/>



Question 121: **Incorrect**

**What are some of the common causes why you cannot connect to a DB instance on AWS ?  
Select 3 options.**

• ☒

A. There is a read replica being created, hence you cannot connect

**(Incorrect)**

• ☒

B. The DB is still being created

**(Correct)**

• ☐

C. The local firewall is stopping the communication traffic

**(Correct)**

• ☒

D. The security groups for the DB are not properly configured.

**(Correct)**

### **Explanation**

There are some steps clearly given in the AWS documentation: For more information on rds troubleshooting please visit the below link

[http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP\\_Troubleshooting.html](http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Troubleshooting.html)

Question 124: **Incorrect**

An image named photo.jpg has been uploaded to a bucket named examplebucket in the us-east-1 region. Which of the below is the right URL to access the image, if it were made public ? Consider that S3 is used as a static website.

☐

A. <http://examplebucket.s3-website-us-east-1.amazonaws.com/photo.jpg>

(Correct)

☐

B. <http://examplebucket.website-us-east-1.amazonaws.com/photo.jpg>

☒

C. <http://examplebucket.s3-us-east-1.amazonaws.com/photo.jpg>

(Incorrect)

☐

D. <http://examplebucket.amazonaws.s3-website-us-east-1./photo.jpg>

### Explanation

The URL for a S3 web site is as shown in the KB Article [.s3-website-.amazonaws.com](#). Hence the right option is option A. When you configure a bucket for website hosting, the website is available via the region-specific website endpoint. Website endpoints are different from the endpoints where you send REST API requests. For more information about the differences between the endpoints, see [Key Differences Between the Amazon Website and the REST API Endpoint](#). The two general forms of an Amazon S3 website endpoint are as follows: --> bucket-name.s3-website-region.amazonaws.com --> bucket-name.s3-website.region.amazonaws.com. Which form is used for the endpoint depends on what Region the bucket is in. For example, if your bucket is named example-bucket and it resides in the US East (N. Virginia) region, the website is available at the following Amazon S3 website endpoint: [For more information on the bucket and the URL format for S3 buckets](#), please visit the below url:

<http://docs.aws.amazon.com/AmazonS3/latest/dev/WebsiteEndpoints.html>

<http://docs.aws.amazon.com/AmazonS3/latest/dev/HostingWebsiteOnS3Setup.html>

Question 125: **Incorrect**

**A company has an EC2 instance that is hosting a web solution which is mostly used for read-only purposes. The CPU utilization is constantly 100% on the EC2 instance. Which of the below solutions can help alleviate and provide a quick resolution to the problem.**

☐

A. Use Cloudfront and place the EC2 instance as the origin

**(Correct)**

☐

B. Let the EC2 instance continue to run at 100%, since the AWS environment can handle the load.

☒

C. Use SNS to notify the IT admin when it reaches 100% so that they can disconnect some sessions to help alleviate the load

**(Incorrect)**

☐

D. Use SES to notify the IT admin when it reaches 100% so that they can disconnect some sessions to help alleviate the load

### **Explanation**

Cloudfront can be used alleviate the load on web based solutions by caching the recent reads in its edge locations and reduce the burden on the EC2 instance. Amazon CloudFront is a global content delivery network (CDN) service that accelerates delivery of your websites, APIs, video content or other web assets For more information on AWS Cloudfront please visit the below url <https://aws.amazon.com/cloudfront/>

Question 127: **Incorrect**

**Which of the below mentioned services are the building blocks for creating a basic high availability architecture in AWS. Select 2 options.**

• ☐

A. EC2

(Correct)

• ☐

B. SQS

• ☒

C. Elastic Load Balancer

(Correct)

• ☐

D. Cloudwatch

### **Explanation**

Having EC2 instances hosting your applications in multiple subnets, hence multiple AZ's and placing them behind an ELB is the basic building block of a high availability architecture in AWS. For more information on High availability and Fault tolerance please visit the below url [https://media.amazonwebservices.com/architecturecenter/AWS\\_ac\\_ra\\_ftha\\_04.pdf](https://media.amazonwebservices.com/architecturecenter/AWS_ac_ra_ftha_04.pdf)

Question 128: **Incorrect**

**You have a set of EC2 Instances launched via Autoscaling. You now want to change the Instance type for the instances that would be launched in the future via Autoscaling. What would you do in such a case**



A. Change the Launch configuration to reflect the new instance type

**(Incorrect)**



B. Change the Autoscaling Group and add the new instance type.



C. Create a new Launch Configuration with the new instance type and replace the existing Launch configuration attached to the Autoscaling Group.

**(Correct)**



D. Create a new Launch Configuration with the new instance type and add it along with the existing Launch configuration attached to the Autoscaling Group.

### **Explanation**

The AWS Documentation mentions the following When you create an Auto Scaling group, you must specify a launch configuration. You can specify your launch configuration with multiple Auto Scaling groups. However, you can only specify one launch configuration for an Auto Scaling group at a time, and you can't modify a launch configuration after you've created it. Therefore, if you want to change the launch configuration for your Auto Scaling group, you must create a launch configuration and then update your Auto Scaling group with the new launch configuration. For more information on Launch Configuration for Autoscaling, please refer to the below link

<http://docs.aws.amazon.com/autoscaling/latest/userguide/LaunchConfiguration.html>

Question 133: **Incorrect**

**You work for a company who are deploying a hybrid cloud approach. Their legacy servers will remain on premise within their own datacenter however they will need to be able to communicate to the AWS environment over a site to site VPN connection. What do you need to do to establish the VPN connection?**



A. Connect to the environment using AWS Direct Connect.

**(Incorrect)**



B. Assign a static routable address to the customer gateway

**(Correct)**



C. Create a dedicated NAT and deploy this to the public subnet.



D. Update your route table to add a route for the NAT to 0.0.0.0/0.

### **Explanation**

This requirement is given in the AWS documentation for the customer gateway. The traffic from the VPC gateway must be able to leave the VPC and traverse through the internet onto the customer gateway. Hence the customer gateway needs to be assigned a static IP that can be routable via the internet. For more information on VPC Virtual Private connections, please refer to the below link: [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_VPN.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_VPN.html)

Question 141: **Incorrect**

**In order to establish a successful site-to-site VPN connection from your on-premise network to the VPC (Virtual Private Cloud), which of the following needs to be configured outside of the VPC? Choose the correct answer from the options below**

☐

A. The main route table to route traffic through a NAT instance

☒

B. A public IP address on the customer gateway for the on-premise network

**(Correct)**

☐

C. A dedicated NAT instance in a public subnet

☐

D. An Elastic IP address to the Virtual Private Gateway

**(Incorrect)**

### **Explanation**

On the customer side gateway you need to have a public IP address which can be addressed by the VPN connection. For more information on VPN connections, please refer to the below link: <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpn-connections.html>

Question 144: **Incorrect**

**Your organization has been using a HSM (Hardware Security Module) for secure key storage. It is only used for generating keys for your EC2 instances. Unfortunately, the HSM has been zeroized after someone attempted to log in as the administrator three times using an invalid password. This means that the encryption keys on it have been wiped. You did not have a copy of the keys stored anywhere else. How can you obtain a new copy of the keys that you had stored on HSM? Choose the correct answer from the options below**

☐

A. You cannot; the keys are lost if you did not have a copy.

**(Correct)**

☒

B. Contact AWS Support; your incident will be routed to the team that supports AWS CloudHSM and a copy of the keys will be sent to you after verification

**(Incorrect)**

☐

C. Restore a snapshot of the HSM

☐

D. You can still connect via CLI; use the command 'get-client-configuration' and you can get a copy of the keys

### **Explanation**

This is given in the aws documentation For more information on CloudHSM, please refer to the below link <https://aws.amazon.com/cloudhsm/faqs/>



Question 146: **Incorrect**

**A customer wants to leverage Amazon Simple Storage Service (S3) and Amazon Glacier as part of their backup and archive infrastructure. The customer plans to use third-party software to support this integration. Which approach will limit the access of the third party software to only the Amazon S3 bucket named “company-backup”?**

☐

A. A custom bucket policy limited to the Amazon S3 API in the Amazon Glacier archive “company-backup”

☒

B. A custom bucket policy limited to the Amazon S3 API in “company-backup”

**(Incorrect)**

☐

C. A custom IAM user policy limited to the Amazon S3 API for the Amazon Glacier archive “company-backup”.

☐

D. A custom IAM user policy limited to the Amazon S3 API in “company-backup”.

**(Correct)**

### **Explanation**

You can use IAM user policies and attach them to users/groups that need specific access to S3 buckets. An example of creating such policies is given in the link below  
<https://aws.amazon.com/blogs/security/writing-iam-policies-how-to-grant-access-to-an-amazon-s3-bucket/>

Retake test

[Continue](#)